



# Welcome to the **SaaS Lab Program**

Session 5

## **Application Storage Modernization**

**This event will be recorded. Your name or other information may end up in the recording. If you do not wish to be recorded, please drop out of this session.**

**The event will start shortly**



**27:00**

# Hello, meet your session presenters



## Daniel Scott-Raynsford

Partner Technology Strategist

About: Daniel is a technology professional with 23 years' experience designing, building, delivering and running applications and systems utilizing a range of software development technologies. Daniel now has a strong focus on cloud architecture, SaaS development and DevOps practices.



[daniel.scottraynsford@microsoft.com](mailto:daniel.scottraynsford@microsoft.com)



<https://www.linkedin.com/in/dscottraynsford/>



If you'd like guidance on your Azure modernization journey, please e-mail the SaaS Lab team

[saaslab@microsoft.com](mailto:saaslab@microsoft.com)

We would love to help!





## Joining AI & ML Session?

Please fill in the survey to help us tailor the session!



<https://forms.office.com/r/hFcUr8XjcY>





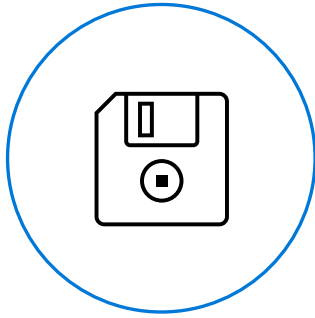
About this session

# Agenda

- Modernizing beyond managed services
- Introduction to storage multi-tenancy
- Break for questions & comments
- A typical storage architecture modernization journey
- Quiz time
- Q&A

# How do applications commonly store data in Azure?

Azure SaaS apps store application data in different services. Here are the [most common\\* ones in Azure](#).



## File Based Storage

- [IaaS Storage Servers](#) (unmanaged)
- [Azure Blob Storage](#)
- [Azure Files](#)



## Relational Databases

- [IaaS Databases](#) (unmanaged)
- [Azure SQL](#)
- [Azure SQL MI](#)
- [Azure Database for PostgreSQL](#)
- [Azure Database for MySQL](#)



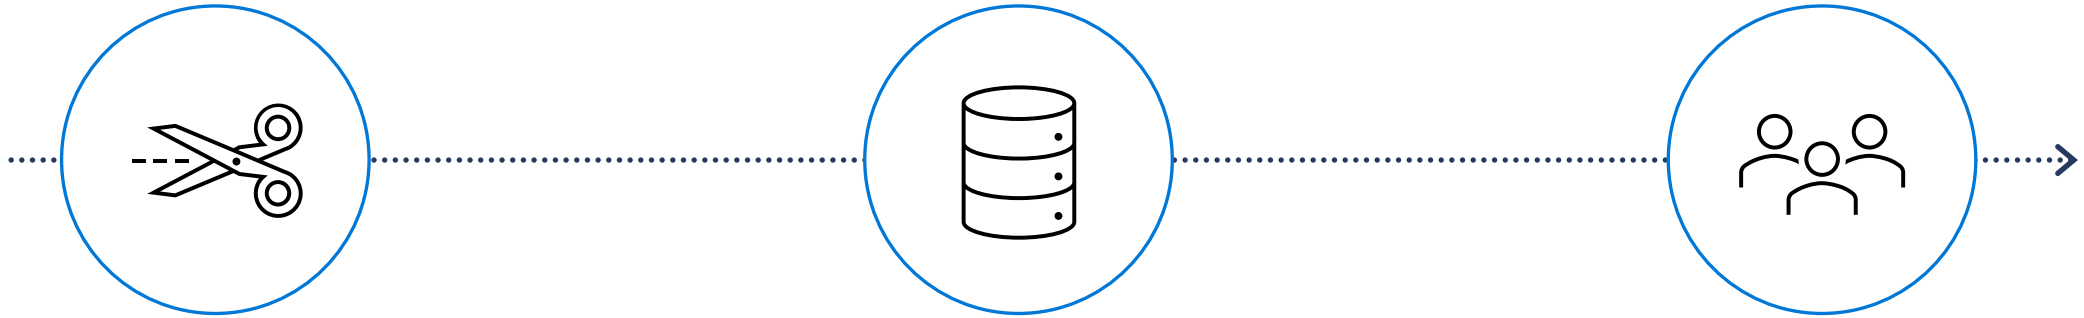
## NoSQL Databases

- IaaS NoSQL (e.g., [MongoDB](#))
- [Azure Table Storage](#)
- [Cosmos DB](#), which includes:
  - Document DB ([MongoDB API](#), [Cassandra API](#), SQL API)
  - Graph DB ([Gremlin API](#))

*\* There are other data store types, such like etcd, but those are out of scope for this session.*

# Storage modernization journey

The common steps a SaaS builder takes to modernize their application persistence layer.



## Decouple Persistence Layer

Separate the application persistence layer (storage) from the application layer.

- **Horizontal scaling of application**
- **Enable high-availability**
- **Modern architecture patterns**

*Decoupling is often combined with moving to managed services.*

## Use Managed Services

Move from unmanaged storage services running in IaaS virtual machines to managed storage services.

- **Reduce operational cost**
- **Increased elasticity**
- **Faster tenant onboarding**

## Multi-tenancy

Enable multi-tenant patterns on storage services.

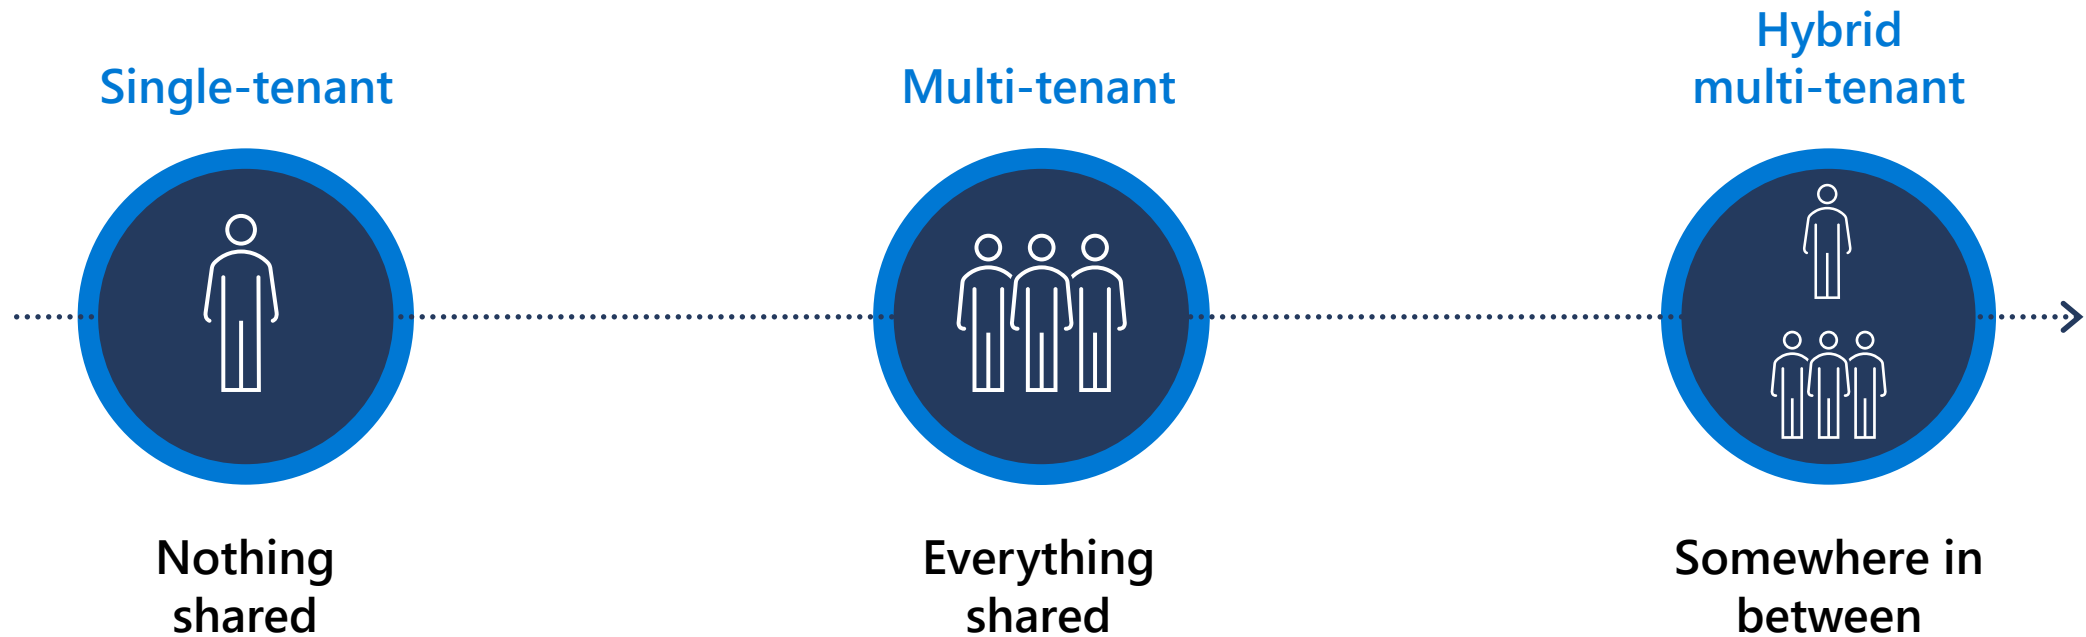
- **Increase resource utilization**
- **Reduce cost per tenant**
- **Increase maximum tenant limits**

*Multi-tenanting of storage can be done with un-managed services, but with reduced effectiveness.*



# SaaS multi-tenancy

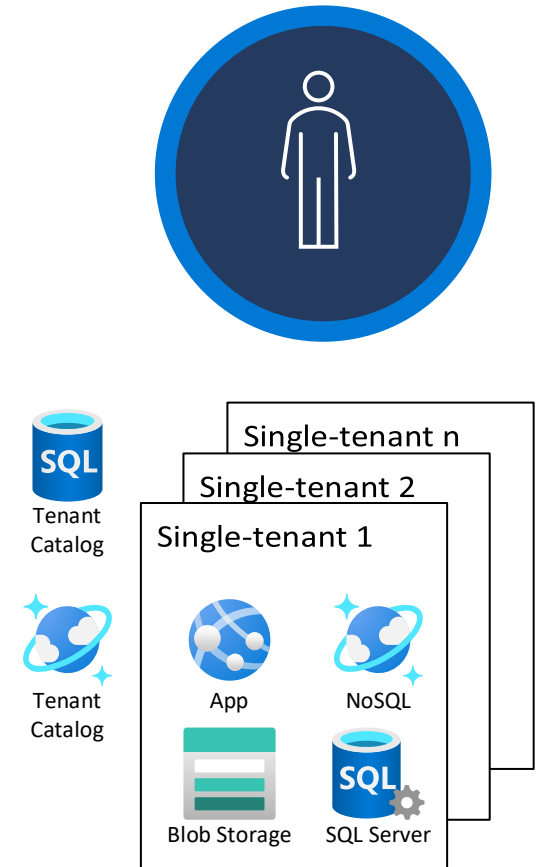
---



# Single-tenant: shared nothing

---

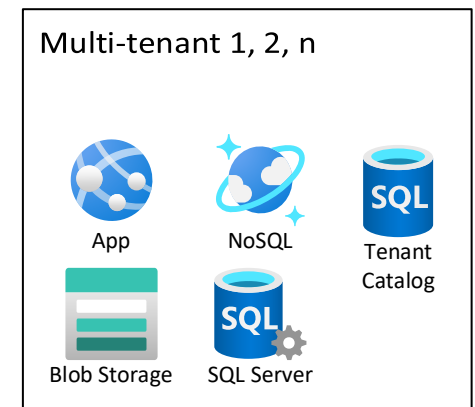
- Single tenant per deployment
  - Storage account per tenant
  - Database per tenant
  - Document/graph store per tenant
- Customization can be done in each deployment
  - ! • This does not scale and will eventually become unmaintainable
- Higher operational costs & more waste
- ☆ • High level of tenant isolation
- ☆ • Easy to understand & measure per-tenant costs



# Multi-tenant: shared everything

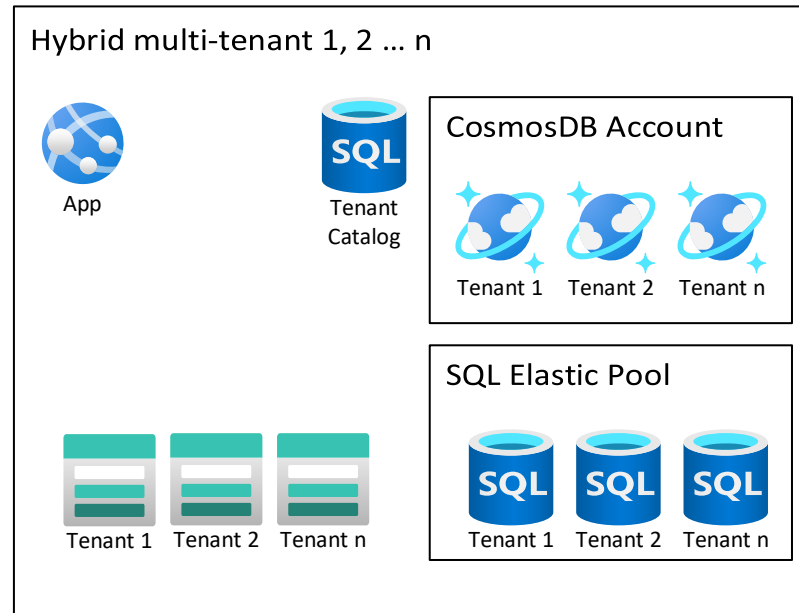
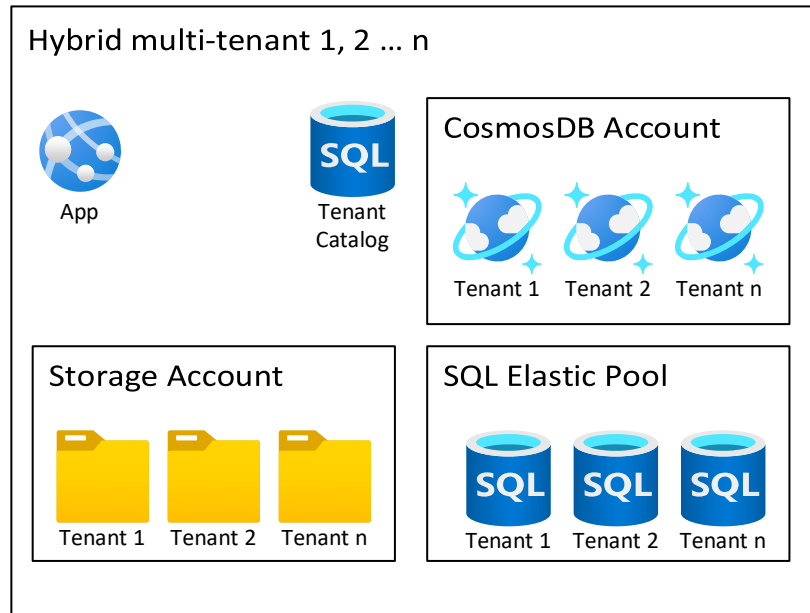
---

- Multiple tenants in a single deployment
  - Single storage account
  - Single database
  - Single document/graph store
- Tenant customization via data
- ! • Lower level of isolation
  - Potential noisy-neighbor
  - Commercial sales challenges
- ! • Difficult to measure per-tenant costs



# Hybrid multi-tenant: somewhere in between

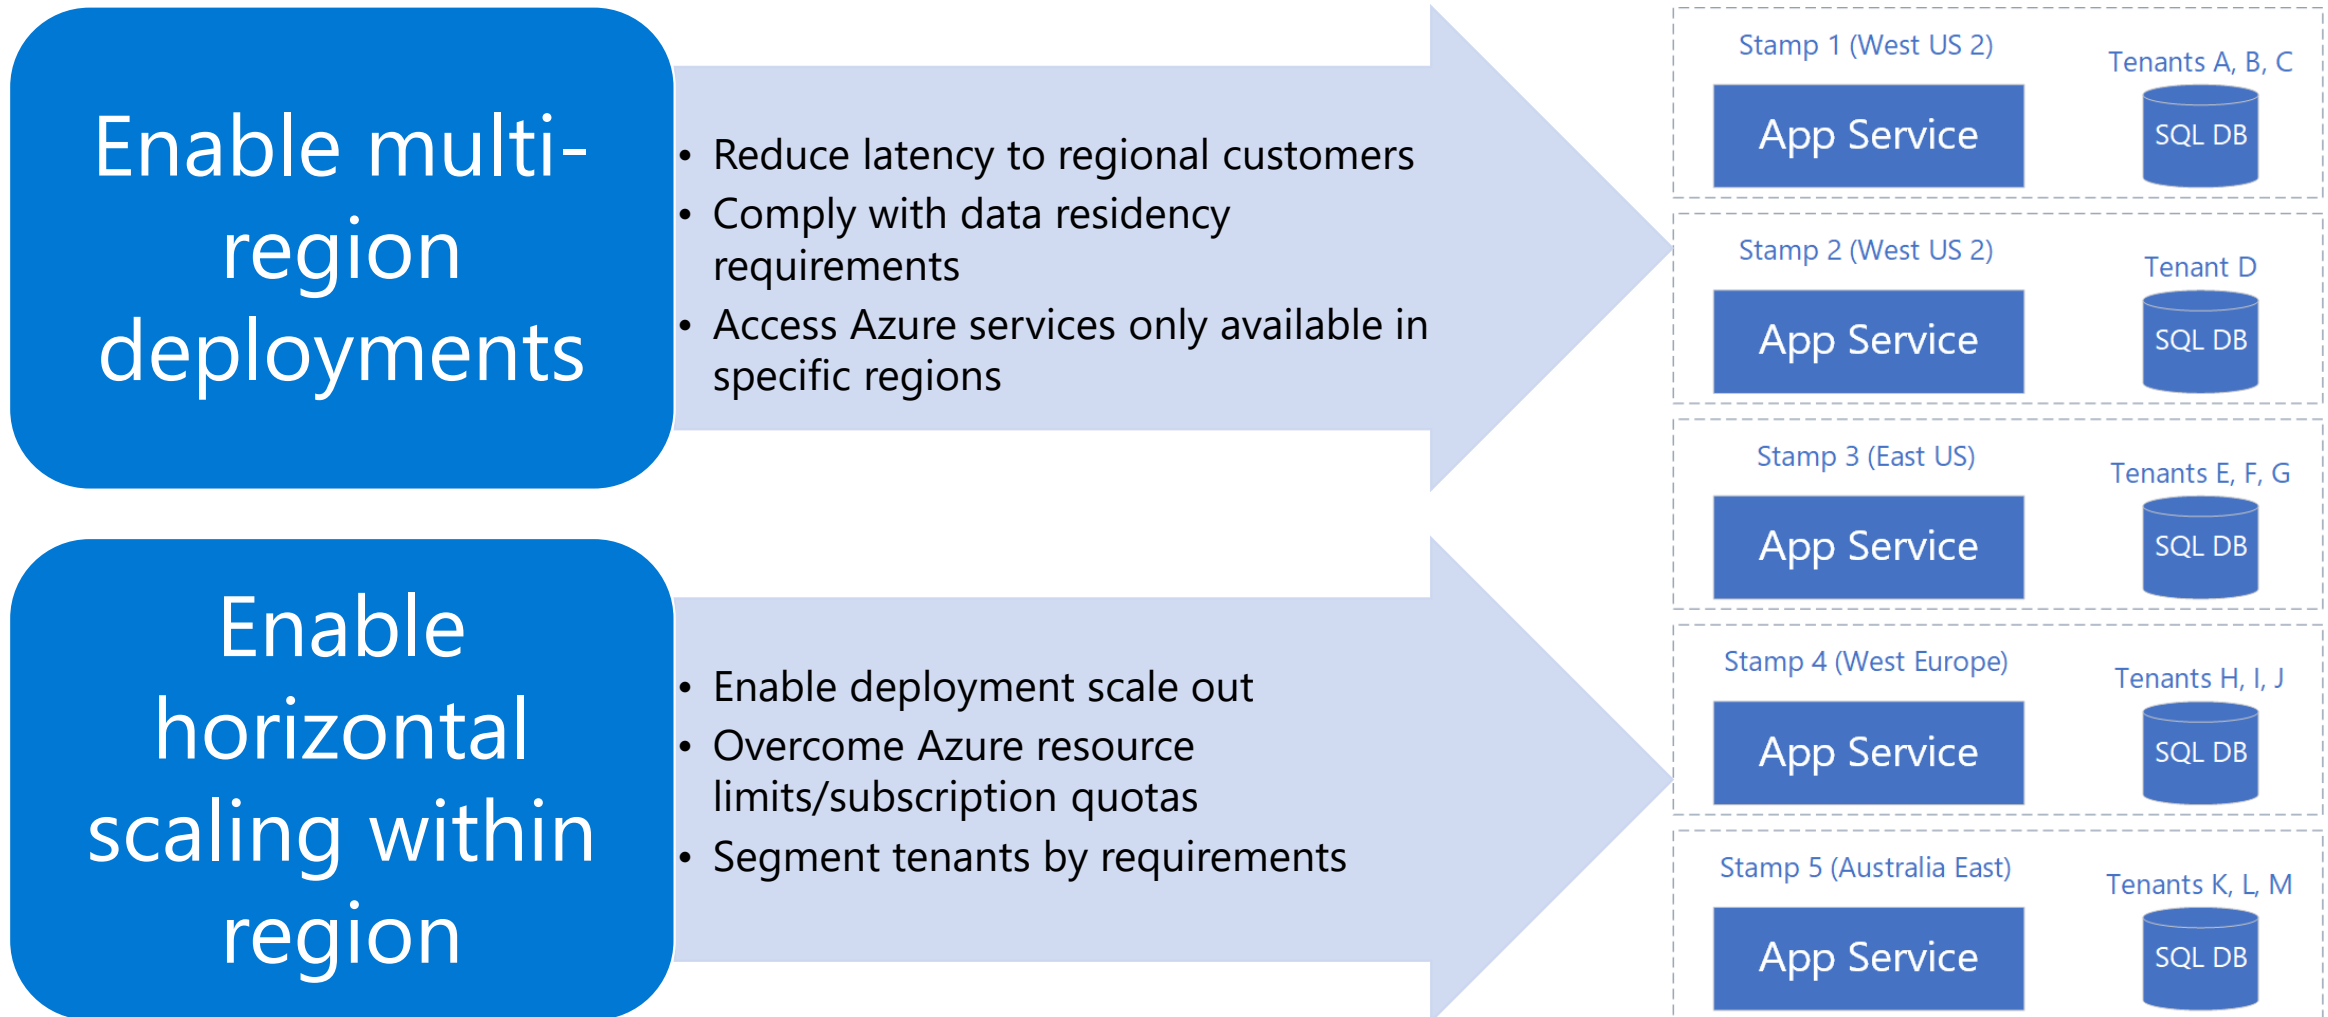
- Many different hybrid models:



- And anywhere in between
- Good balance between single & multi-tenant
- Enable different commercial models




# Deployment stamps architecture



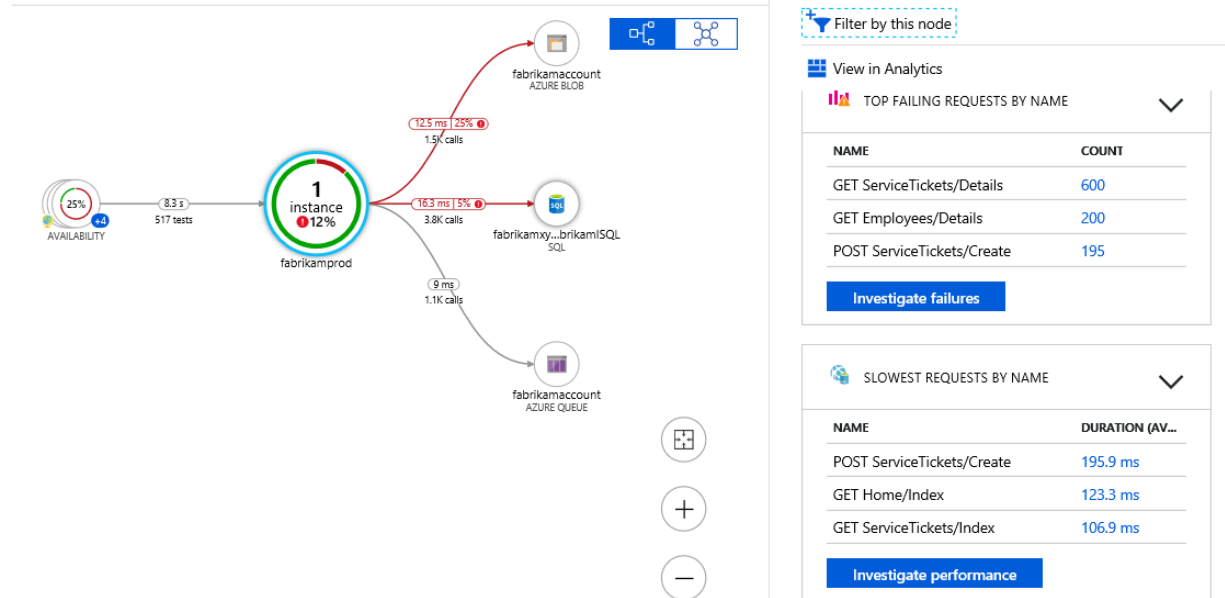


Where are you on your SaaS application storage architecture journey - choose the closest model?

 Start presenting to display the poll results on this slide.

# Monitoring multi-tenant storage

- Enable [Application Insights](#) on your application
- Identify **key tenant storage operations** and **continuously monitor baseline performance**
- Enable [Application Insights Smart Detection](#) - Watch out for [performance anomalies](#)
- [Send custom telemetry](#) w/ **tenant information** to Application Insights
- Use [ITelemetryInitializer](#) to enrich metrics with TenantID.
- Set up alerts to notify of tenant or system performance degradation
- Trigger automation to mitigate performance conditions (e.g., migrate tenant)
- Use [Application Map](#) to trace bottlenecks and tenant performance issues



# Application Insights

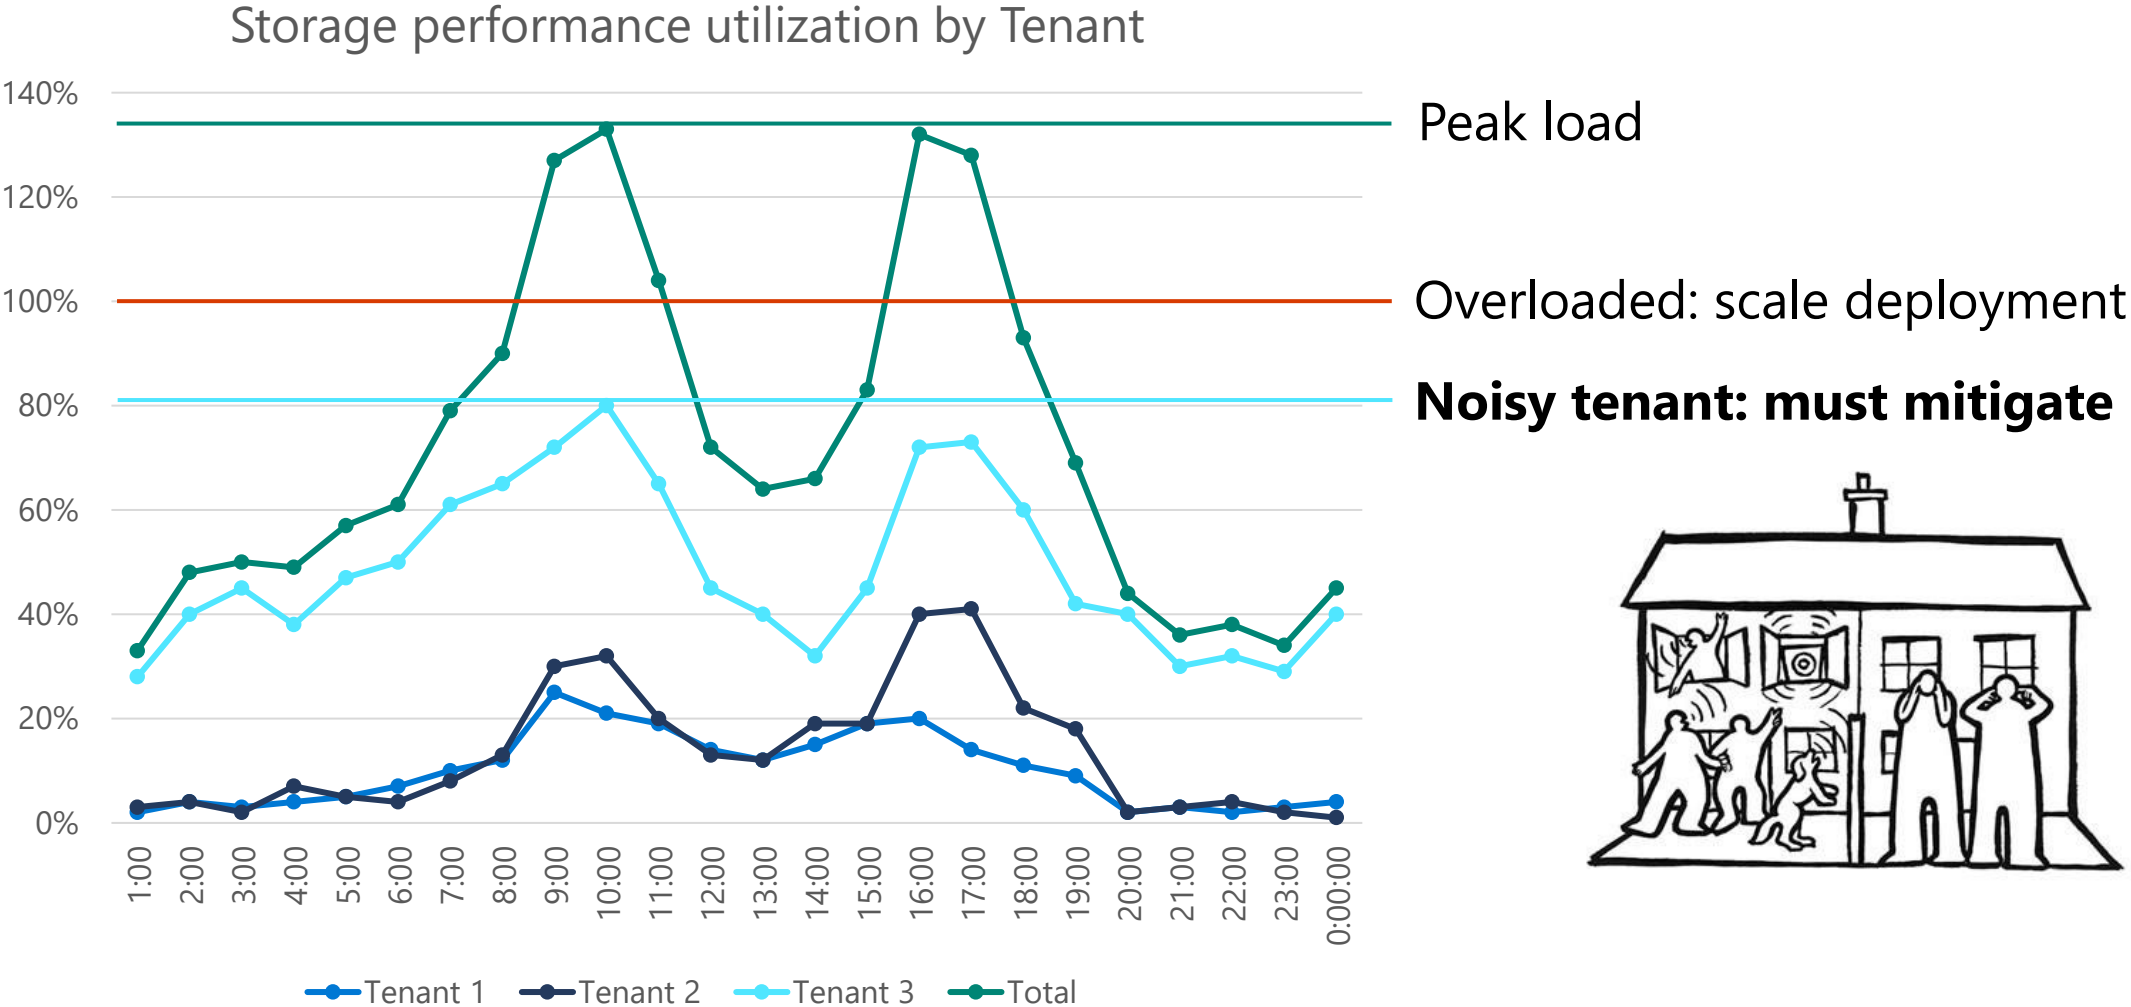
A quick look at [Application Insights](#)

Adding tenant information with [ITelemetryInitializer](#)

Monitoring with [Application Map](#)



# Deployment & tenant storage utilization



# Mitigating noisy tenants

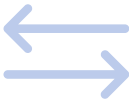
---



Enable auto-scale on Azure services



Optimise/modernize storage services



Migrate to deployment stamp with lower load



Migrate to new deployment stamp




Cordon the tenant with request/usage throttling



Consider requiring purchase of a higher service tier





Let's have a  
break for a  
moment...

...do you have  
any questions  
or comments?





A typical storage architecture modernization journey

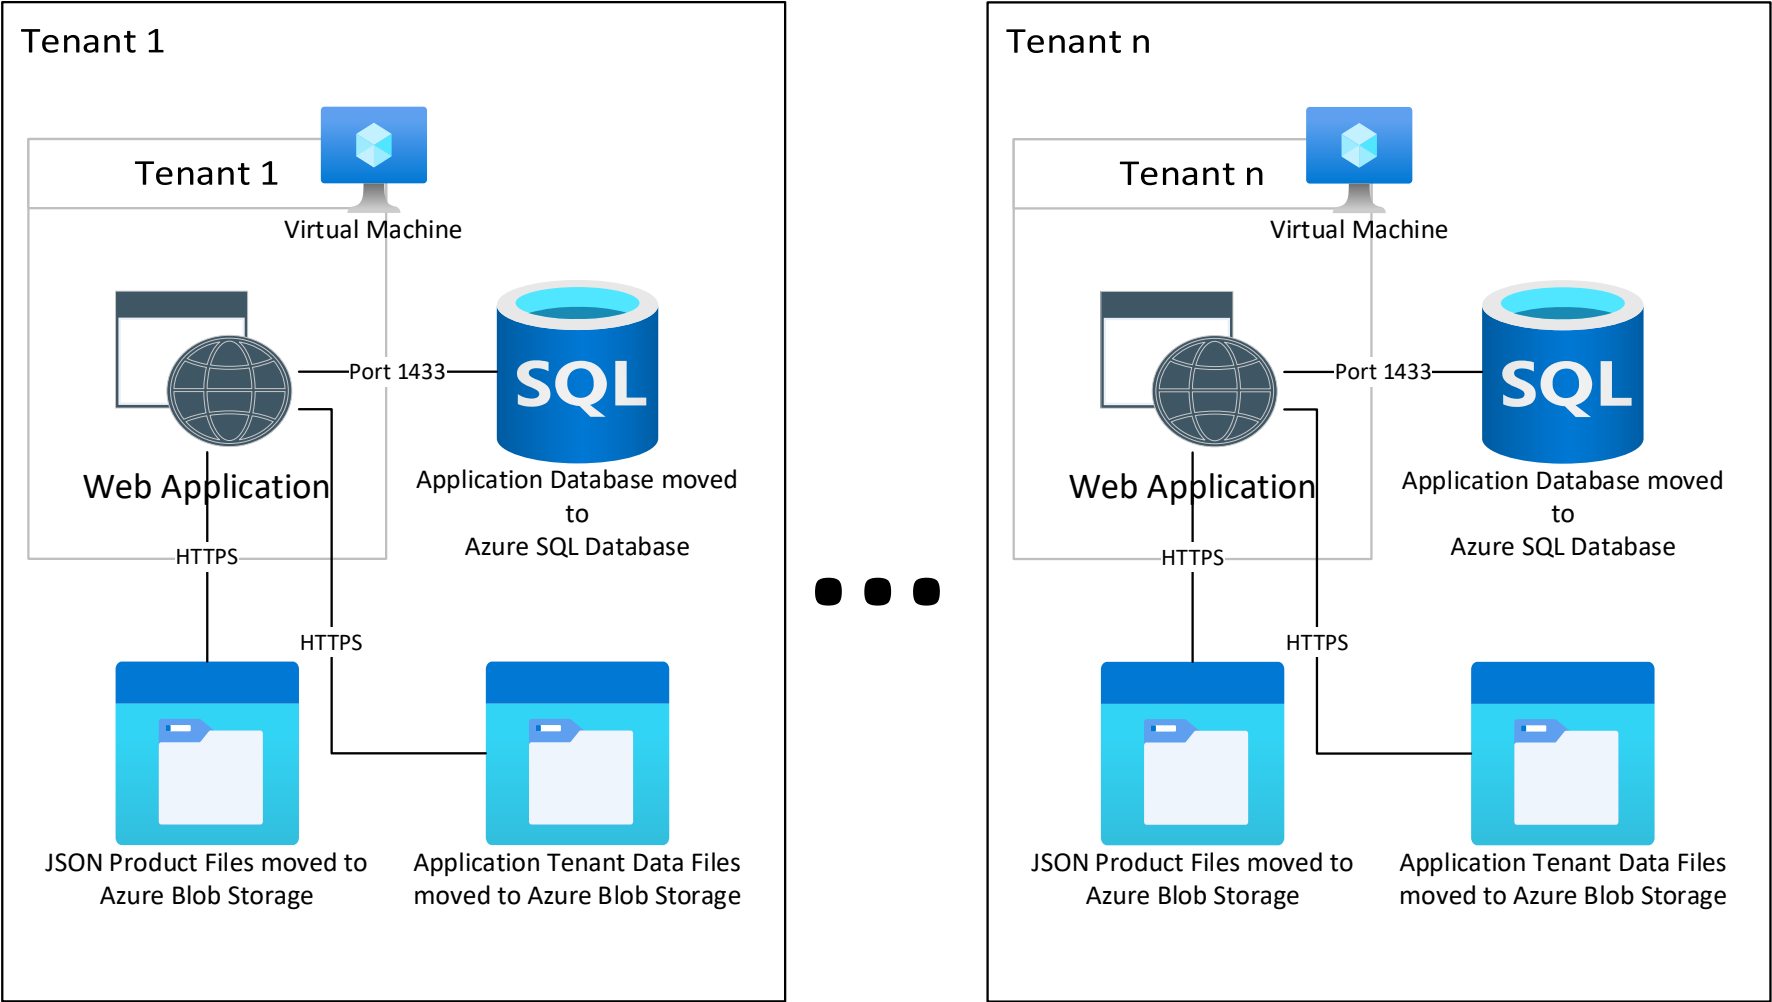




- Recently migrated an on-premises version of their retail services solution as a **SaaS application in Azure**.
- Are expecting to onboard at least **5000 tenants** in 3 years.
- Plan to enable new commercial models like **trial** or **monthly offers, paid add-ons** and a **premium tier**.
- Need to **reduce operational cost per tenant** and **onboard time**.
- Delight customers with **high-quality, reliable service** and continuous delivery of **innovative new features**.
- A managed identity platform, Azure AD, is used.



# Initial single-tenant architecture

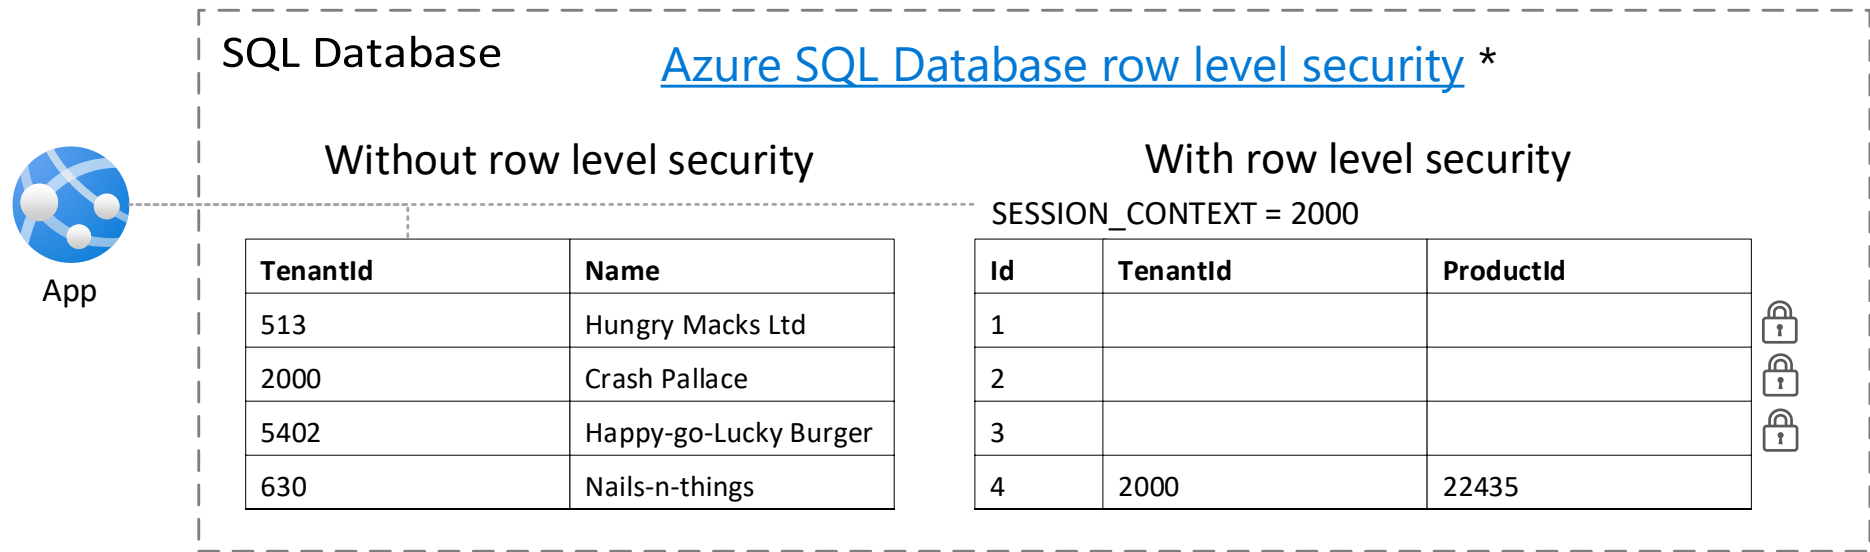






# Azure SQL multi-tenant patterns

# Table multi-tenancy

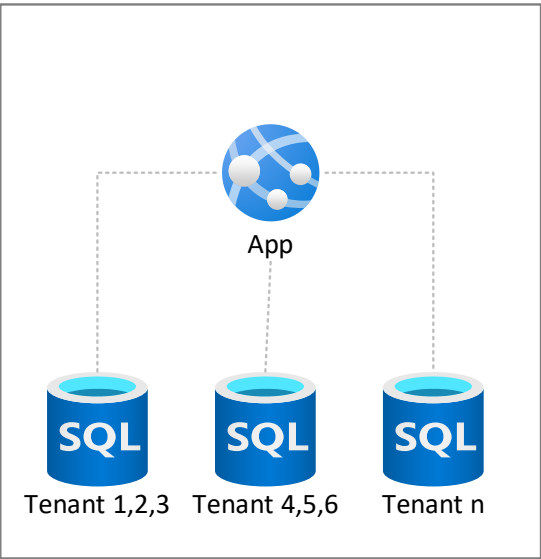
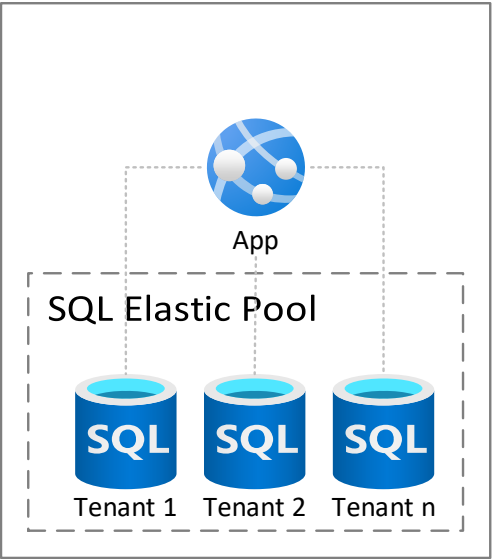
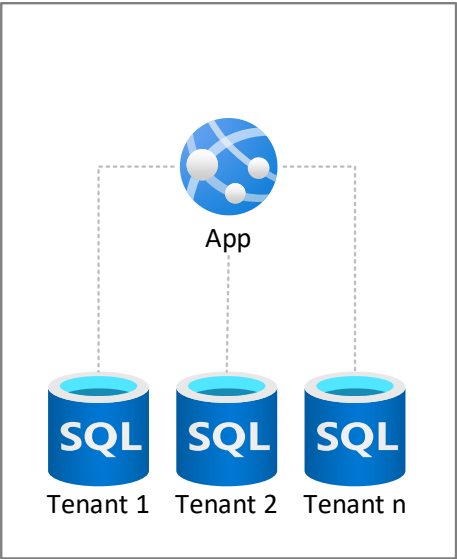


	Without row level security	With row level security
Scale	Low, 1-1,000s	Low, 1-1,000s
Tenant isolation	Low	Medium
Database cost per tenant	Lowest	Lowest
Development complexity	Medium	Medium
Operational complexity	Low	Low

\* Row-level security is supported in other RDBMS in Azure but is not covered here.



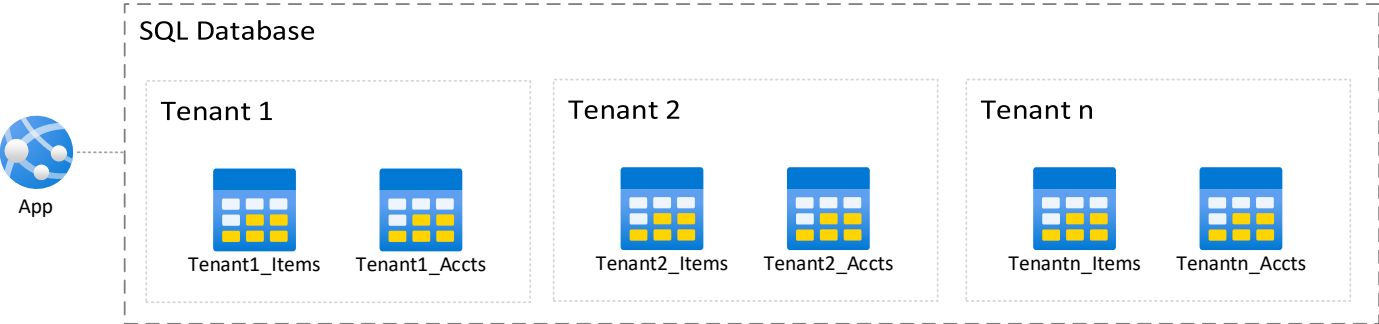
# Azure SQL database patterns



	Database per-tenant	Elastic pool database-per-tenant	Sharded multi-tenant databases
Scale	Very High, 1-100,000s	Very High, 1-100,000s	Unlimited, 1-1,000,000s
Tenant isolation	High	High	Low except for singleton tenant in a database
Database cost per tenant	Low	Low	Lowest
Development complexity	Low	Low	Medium due to sharding
Operational complexity	Low-Medium	Low-Medium	Low-High

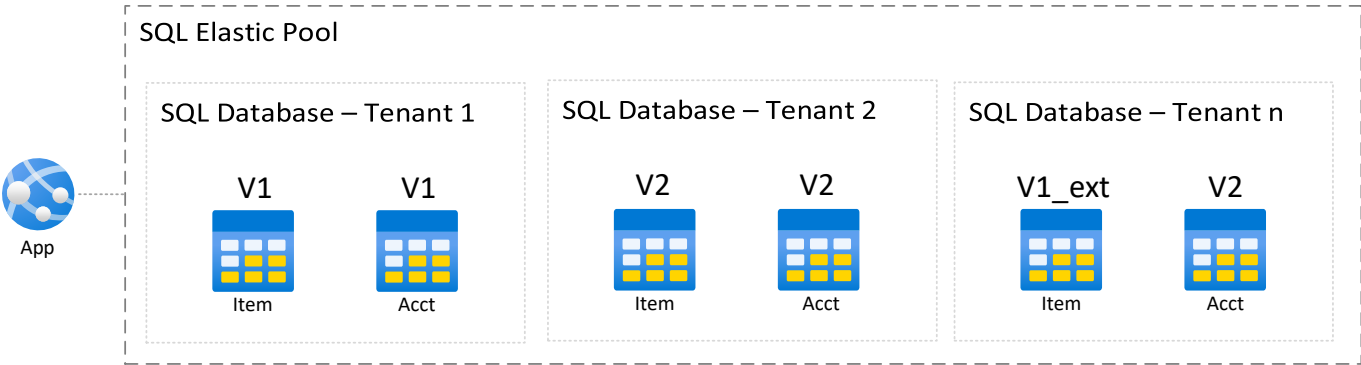
# SQL multi-tenant anti-patterns

## ❌ Single database with multi-tenant table set



Scale	Very Low, 1-100s
Tenant isolation	Low
Database cost per tenant	Low
Development complexity	High; increasing with scale
Operational complexity	High; increasing with scale

## ❌ Per-tenant schema customizations



Scale	Very High, 1-100,000s
Tenant isolation	High
Database cost per tenant	Low
Development complexity	High; increasing with scale
Operational complexity	High; increasing with scale

Consider adopting NoSQL if per tenant schema customization is required.

# Consider Azure limits

---

## Azure SQL Server limits

As your application scales to 100s, 10,000s or even 1,000,000s will consider how cost per tenant and performance will be impacted.

Databases per server 5000

Default number of servers per subscription in any region 50

Max number of servers per subscription in any region 200

DTU / eDTU quota per server 54,000

vCore quota per server instance 140

Max pools per server Limited by number of DTUs or vCores. For example, if each pool is 100 DTU, then server can support 54 pools.

## Azure SQL Elastic Pool limit

Q: How will large numbers of connections affect your front end?

Q: Will you have issues with connection exhaustion?

Q: Can one tenant disrupt service for another?

Consider all pillars of the [well-architected framework](#).

[Resource limits for Azure SQL Database](#)

[SQL Database resource limits for single databases](#)

[SQL Database resource limits for elastic pools and pooled databases](#)


[SQL Database resource limits for SQL Managed Instance](#)

[Limitations in Azure Database for MySQL](#)

[Limitations in Azure Database for PostgreSQL](#)

# Azure SQL multi-tenant

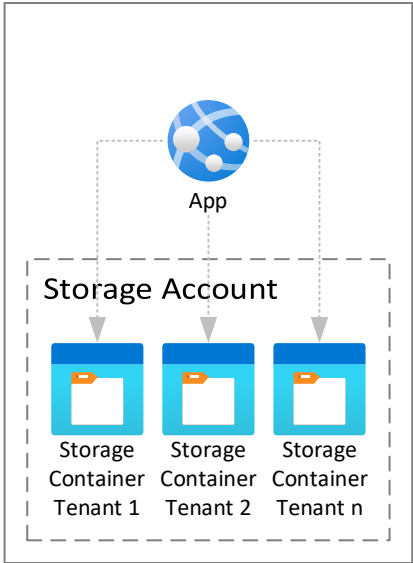
A brief look at a multi-tenanted Azure SQL Databases



# Azure Blob Storage multi-tenant patterns

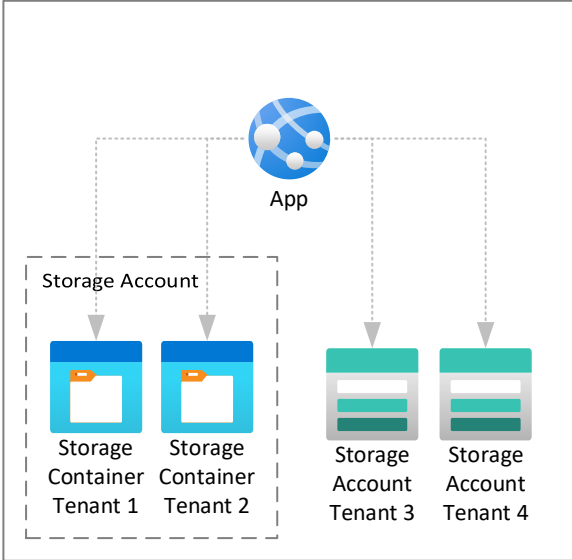
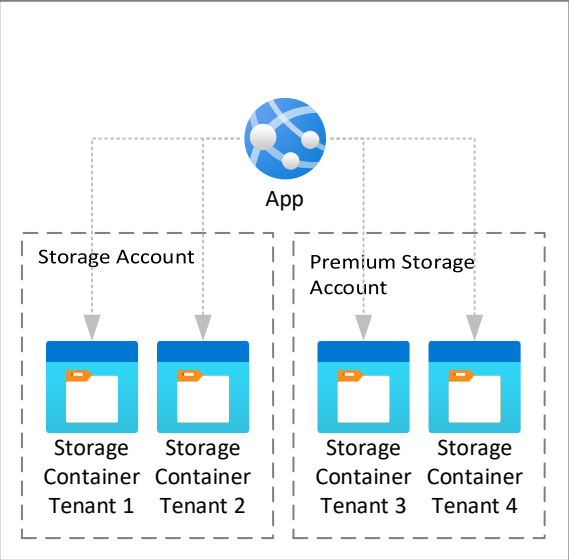


# Azure Blob Storage multi-tenant patterns



	Account per-tenant	Shared account with container per-tenant
Scale	Very Low, 1-250	Unlimited, 1-1,000,000s
Tenant isolation	High	Low
Performance monitoring and management	Per-tenant	Per-account, per-tenant difficult
Development complexity	Low	Low
Operational complexity	Low	Low

# Azure Blob Storage hybrid multi-tenant patterns



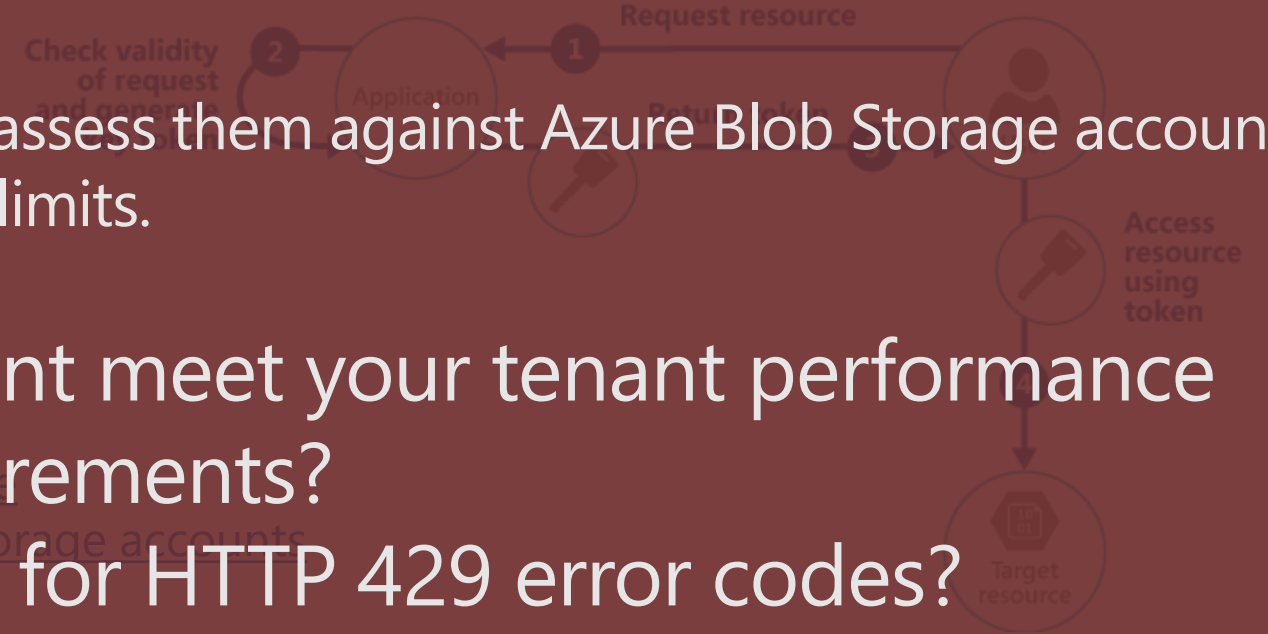
	Shared account with container per-tenant with optional tenant service tiers	Shared account with container per-tenant with account per-tenant option
Scale	Unlimited, 1-1,000,000s	Unlimited, 1-1,000,000s
Tenant isolation	High	Low/High
Performance monitoring and management	Per-account, per-tenant difficult	Per-tenant + Per-account, per-tenant difficult
Development complexity	Low-Medium	Low-Medium
Operational complexity	Low-Medium	Low-Medium

# Considerations

## Azure Storage Account limits

Resource	Limit
Storage Accounts per region per subscription	250
Maximum storage account capacity	5 PiB
Maximum request rate per storage account	20,000 requests per second

## Consider using valet key pattern



Consider blob tenant access patterns and assess them against Azure Blob Storage account limits.

Q: Will a single storage account meet your tenant performance requirements?

Q: Are you monitoring for HTTP 429 error codes?

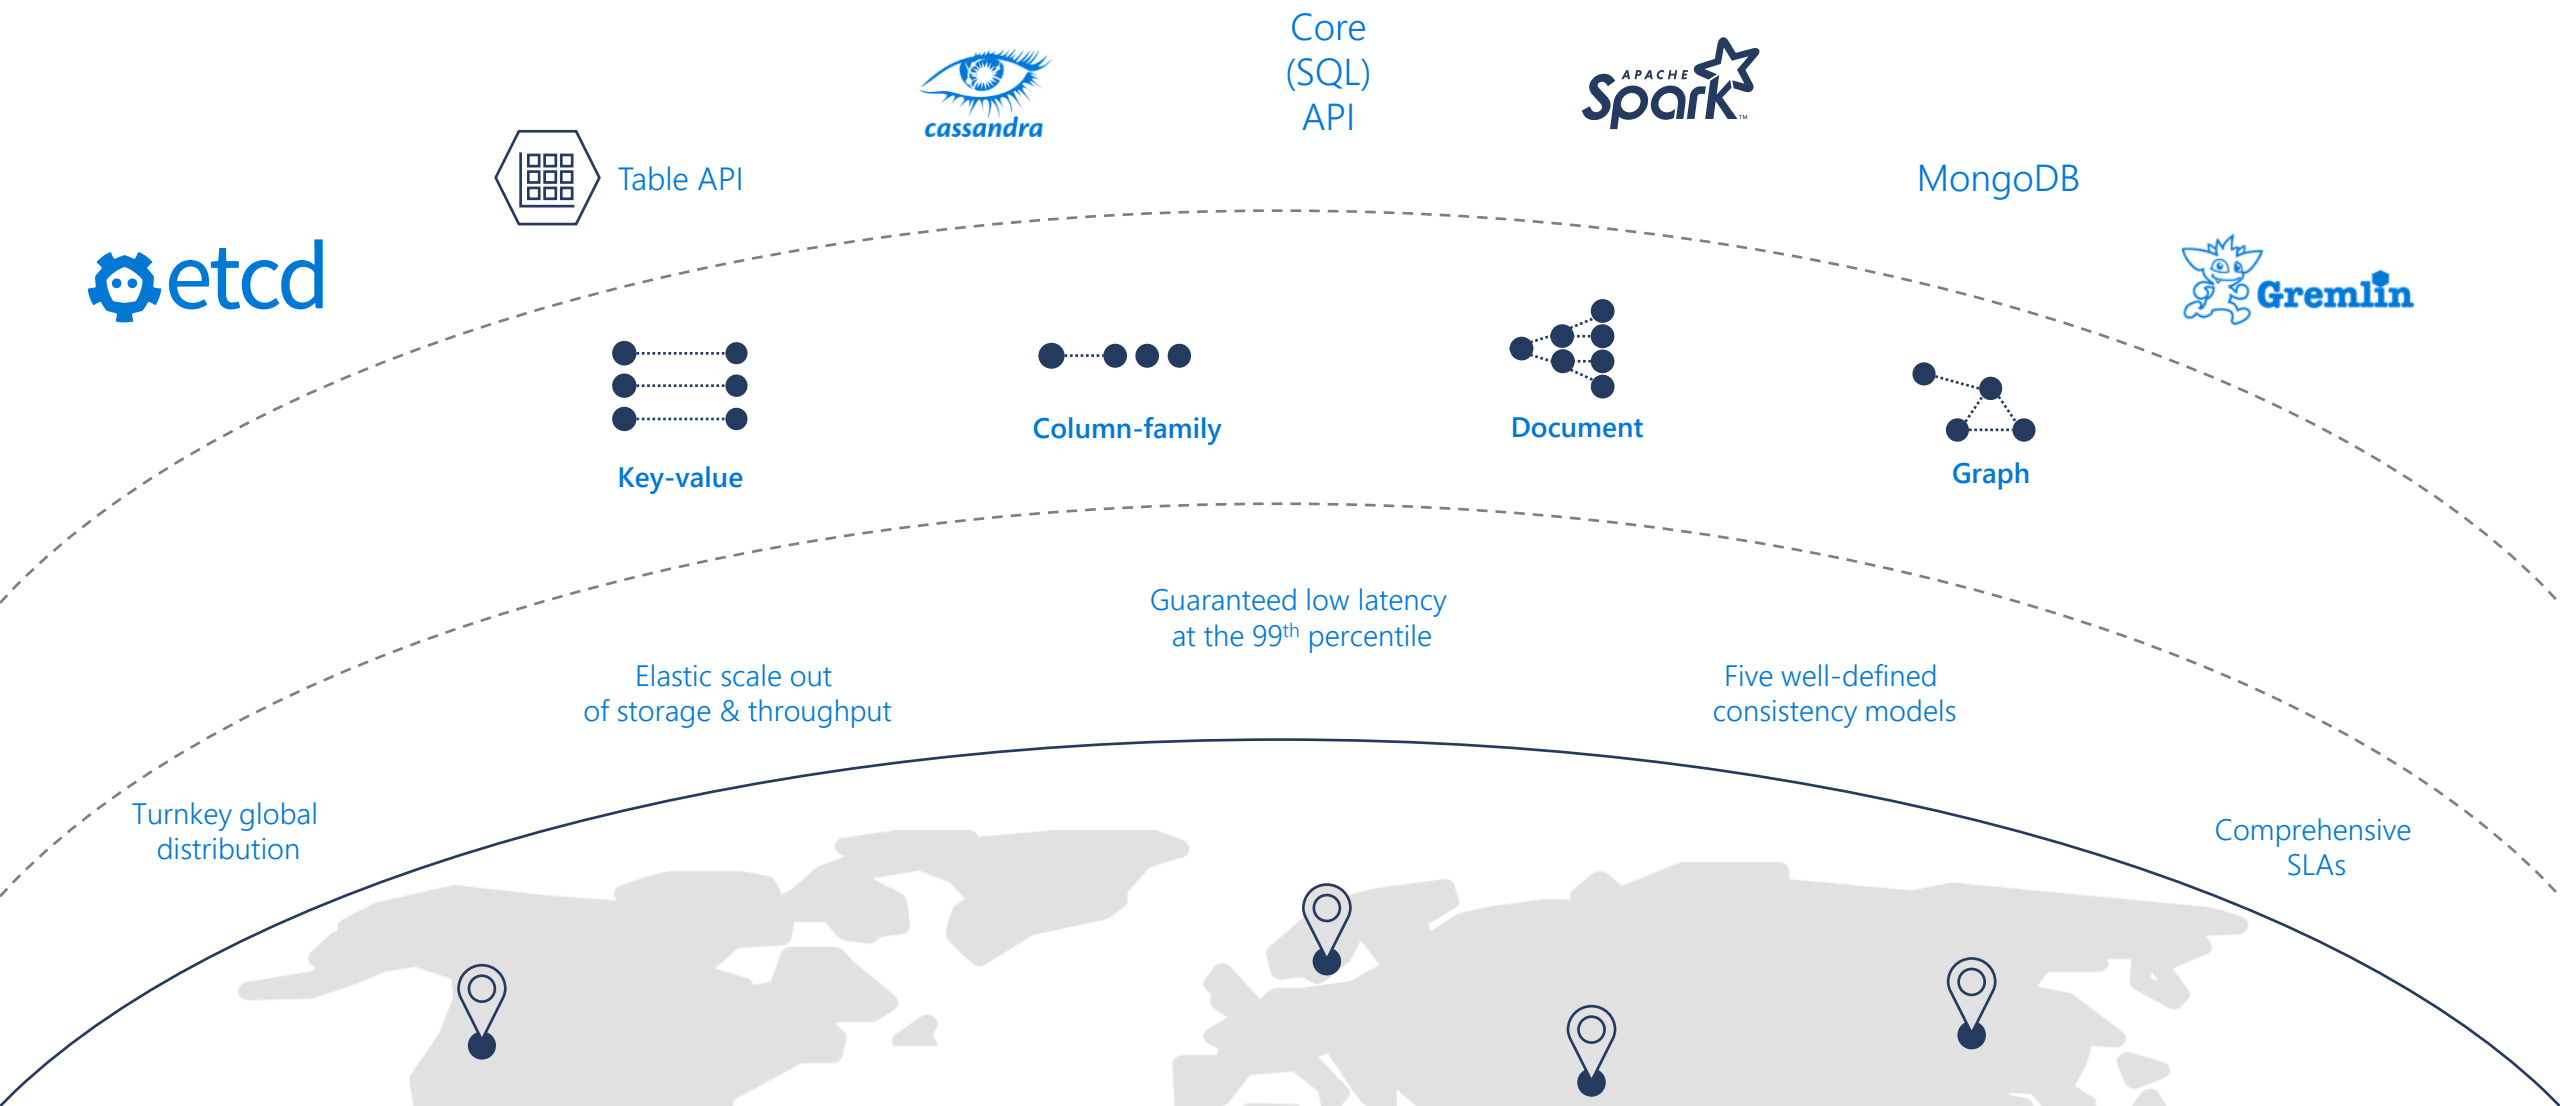
- Model blob tenant access patterns and assess against limits:
  - Ingress/Egress
  - Requests per second
- Clients will receive 429 errors if throttled.
  - SDKs will usually automatically retry, but performance will degrade.

# Azure Blob Storage multi-tenant

A brief look at a multi-tenanted Azure Blob Storage containers

# Azure Cosmos DB

A globally distributed, massively scalable, multi-model database service



# Cosmos DB advantages for multi-tenant

---

- Different tenancy models
- Modern cloud design patterns
- [Dedicated](#) and shared throughput on databases & containers
- Supports [autoscale throughput](#)
- Scale across regions enables [Geode architecture pattern](#) \*
- Easy to measure tenant consumption

# Cosmos DB multi-tenant patterns

	Database Account (per tenant)	Container w/ Dedicated Throughput (per tenant)	Container w/ Shared Throughput (per tenant)	Partition Key (per tenant)
<b>Isolation Knobs</b>	<p>Independent geo-replication knobs</p> <p>Multiple throughput knobs (dedicated throughput – eliminating noisy neighbors)</p>	<p>Independent throughput knobs (dedicated throughput – eliminating noisy neighbors)</p> <p>Group tenants within database account(s) based on regional needs</p>	<p>Share throughput across tenants grouped by database (great for lowering cost on “spiky” tenants)</p> <p>Easy management of tenants (drop container when tenant leaves)</p> <p>Mitigate noisy-neighbor blast radius (group tenants by database)</p>	<p>Share throughput across tenants grouped by container (great for lowering cost on “spiky” tenants)</p> <p>Enables easy queries across tenants (containers act as boundary for queries)</p> <p>Mitigate noisy-neighbor blast radius (group tenants by container)</p>
<b>Throughput requirements</b>	>400 RUs per Tenant (> \$24 per tenant)	>400 RUs per Tenant (> \$24 per tenant)	> 100 RUs per Tenant (> \$6 per tenant)	>0 RUs per Tenant (> \$0 per tenant)
<b>T-Shirt Size</b>	<p>Large</p> <p>Example: Premium offer for B2B apps</p>	<p>Large</p> <p>Example: Premium offer for B2B apps</p>	<p>Medium</p> <p>Example: Standard offer for B2B apps</p>	<p>Small</p> <p>Example: B2C apps</p>



# Considerations

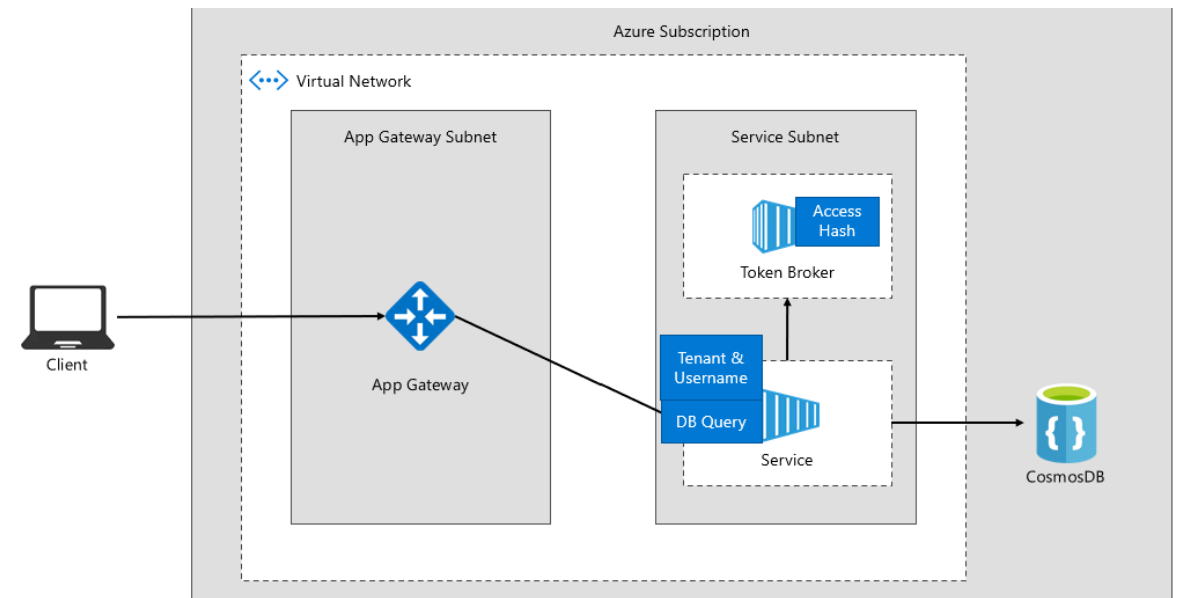
## Provisioned throughput limits

Resource	Limit
Maximum number of databases	Unlimited
Maximum number of containers per database with shared throughput	25
Maximum number of containers per database or account with dedicated throughput	Unlimited

## [Azure Cosmos DB service quotas](#)

- Partition key modelling is critical
- Clients will receive 429 errors if throttled.
  - SDKs will usually automatically retry, but performance will degrade.

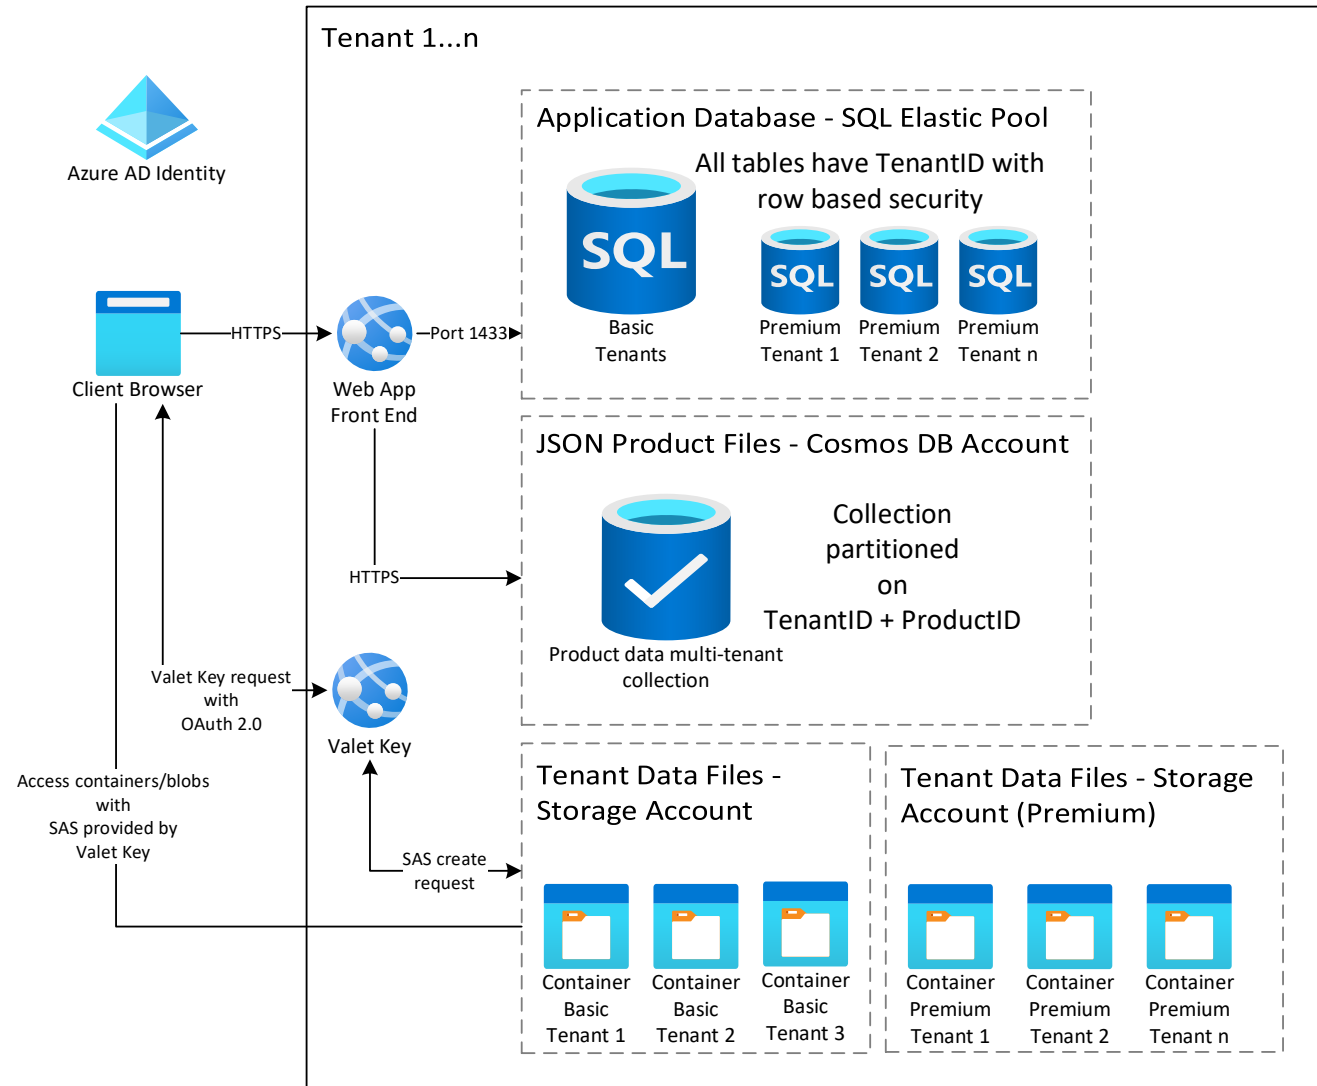
Use [Token Broker pattern](#) for tenant security



# Azure Cosmos DB multi-tenant

A brief look at a multi-tenanted Azure Cosmos DB

# Multi-tenant with architecture



What else should we consider in more detail?

- [Reliability](#)
- [Security](#)
- Limitations of front-end services

# Knowledge check time!



<http://www.kahoot.it> or use the Kahoot! mobile app

<Game Pin>

# Session takeaways

## Continuously Improve

Storage modernization is a journey. Continuously evaluate your customers needs.

Evaluate and review new storage technologies and architectural methodologies for benefits.

Don't be afraid to **rearchitect**.

## Plan for Future Scale

Architect your solution to account for future customer growth, selling into different regions and offering different commercial models.

## Adopt a strategic Multi-tenant approach

Plan your multi-tenant or hybrid multi-tenant architecture carefully, considering:

- Customer needs
- Commercial models
- Azure service limits
- Development and Operational complexity

## Avoid storage anti-patterns

There are patterns that will inhibit your future growth and result in an unmanageable, unscalable or wasteful solution. Avoid these.

## Monitor & Protect

Multi-tenant carefully by monitoring and protecting your customers. Use Application Insights and custom telemetry.

Monitor key tenant baseline metrics such as page load time for important pages.

## Use features of managed services

Managed storage services in Azure provide many features that can help you reduce costs and provide new features for your customers.

Consider additional value you could add just by using the services provided.





Q & A





# Your feedback is important

Please help us improve this program by completing this short feedback form.



<https://aka.ms/saaslabfeedback5>



If you'd like more help on your Azure modernization journey, please e-mail the SaaS Lab team

[saaslab@microsoft.com](mailto:saaslab@microsoft.com)

Thank you for being part of the SaaS Lab Program